



1 Domain Finetuning

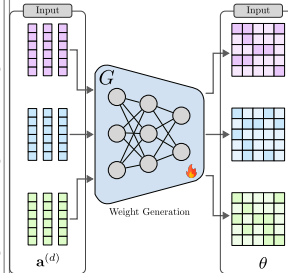


2 Safety Heads Training



- f_θ is a lightweight side network attached to the Finetuned LLM backbone
- θ denotes the SafeHead parameters.
- Finetuned LLM is always frozen.

3 Hypernetwork Training



- Training Losses
- | Weight Recon. | Classifier | Behavior |
|------------------------------|----------------------------|-----------------------------|
| $\mathcal{L}_{\text{recon}}$ | $\mathcal{L}_{\text{cls}}$ | $\mathcal{L}_{\text{func}}$ |

4 Deployment

